



```
%tensorflow_version 2.x
import pandas as pd
import tensorflow as tf
from matplotlib import pyplot as plt
import os
from google.colab import files
uploaded = files.upload()

from numpy import argmax
import re
import numpy as np
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.
Saving Seasons_Stats.csv to Seasons_Stats.csv

```
training_df = pd.read_csv("Seasons_Stats.csv")
```

```
# Print the first rows of the pandas DataFrame.
training_df.head()
```

 **Unnamed: 0** **Year** **Player** **Pos** **Age** **Tm** **G** **GS** **MP** **PER** **TS%** **3PAr** **FTr**

| | | | | | | | | | | | | | |
|----------|---|--------|-----------------|-----|------|-----|------|-----|-----|-----|-------|-----|-------|
| 0 | 0 | 1950.0 | Curly Armstrong | G-F | 31.0 | FTW | 63.0 | NaN | NaN | NaN | 0.368 | NaN | 0.467 |
| 1 | 1 | 1950.0 | Cliff Barker | SG | 29.0 | INO | 49.0 | NaN | NaN | NaN | 0.435 | NaN | 0.387 |
| 2 | 2 | 1950.0 | Leo Barnhorst | SF | 25.0 | CHS | 67.0 | NaN | NaN | NaN | 0.394 | NaN | 0.259 |
| 3 | 3 | 1950.0 | Ed Bartels | F | 24.0 | TOT | 15.0 | NaN | NaN | NaN | 0.312 | NaN | 0.395 |
| 4 | 4 | 1950.0 | Ed Bartels | F | 24.0 | DNN | 13.0 | NaN | NaN | NaN | 0.308 | NaN | 0.378 |

```
training_df["PTS"] /= 1000.0   # Scale the label.
```

```
pd.options.display.max_rows = 10
pd.options.display.float_format = "{:.1f}".format
```

```
test_df = pd.read_csv("Seasons_Stats.csv")
```

```
test_df=test_df[test_df.PTS > 50]
test_df=test_df[test_df.Year >= 2016]    #Test for years since 2016
test_df.head()
```

```
training_df = pd.read_csv("Seasons_Stats.csv")
training_df=training_df[training_df.Year < 2016]
training_df=training_df[training_df.Year > 2012]
```

```
training_df.head()
```

```
training_df[training_df.PTS > 50]
```

↗

| | Unnamed: 0 | Year | Player | Pos | Age | Tm | G | GS | MP | PER | TS% | 3PAr | F |
|--------------|------------|--------|-------------------|-----|------|-----|------|------|--------|------|-----|------|---|
| 21679 | 21679 | 2013.0 | Quincy Acy | PF | 22.0 | TOR | 29.0 | 0.0 | 342.0 | 15.9 | 0.6 | 0.0 | C |
| 21680 | 21680 | 2013.0 | Jeff Adrien | PF | 26.0 | CHA | 52.0 | 5.0 | 713.0 | 13.4 | 0.5 | 0.0 | C |
| 21681 | 21681 | 2013.0 | Arron Afflalo | SF | 27.0 | ORL | 64.0 | 64.0 | 2307.0 | 13.0 | 0.5 | 0.3 | C |
| 21683 | 21683 | 2013.0 | Cole Aldrich | C | 24.0 | TOT | 45.0 | 0.0 | 388.0 | 11.1 | 0.6 | 0.0 | C |
| 21686 | 21686 | 2013.0 | LaMarcus Aldridge | PF | 27.0 | POR | 74.0 | 74.0 | 2790.0 | 20.4 | 0.5 | 0.0 | C |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 23511 | 23511 | 2015.0 | Thaddeus Young | PF | 26.0 | TOT | 76.0 | 68.0 | 2434.0 | 15.7 | 0.5 | 0.1 | C |
| 23512 | 23512 | 2015.0 | Thaddeus Young | PF | 26.0 | MIN | 48.0 | 48.0 | 1605.0 | 15.0 | 0.5 | 0.1 | C |
| 23513 | 23513 | 2015.0 | Thaddeus Young | PF | 26.0 | BRK | 28.0 | 20.0 | 829.0 | 17.1 | 0.5 | 0.2 | C |
| 23514 | 23514 | 2015.0 | Cody Zeller | C | 22.0 | CHO | 62.0 | 45.0 | 1487.0 | 14.1 | 0.5 | 0.0 | C |
| 23515 | 23515 | 2015.0 | Tyler Zeller | C | 25.0 | BOS | 82.0 | 59.0 | 1731.0 | 18.9 | 0.6 | 0.0 | C |

1489 rows × 53 columns

```
del training_df['Year']
```

```
del training_df['Unnamed: 0']
del training_df['Player']
del training_df['Pos']
del training_df['Tm']
del training_df['Age']
del training_df['blan1']
del training_df['blank2']
```

```
del test_df['Year']
del test_df['Unnamed: 0']
del test_df['Player']
del test_df['Pos']
del test_df['Tm']
del test_df['Age']
del test_df['blan1']
del test_df['blank2']
```

#inefficient way to delete all non numeric variables from test and training data fr
#For now I don't want to need at player names, but in the future
#I would use an integer enocding to learn more about player specific performance

```
# Normalize
train_df_mean = training_df.mean()
train_df_std = training_df.std()
train_df_norm = (training_df - train_df_mean)/train_df_std
```

```
test_df_mean = test_df.mean()
test_df_std = test_df.std()
test_df_norm = (test_df - test_df_mean)/test_df_std
```

```
test_df.head()
training_df.head()
```



| | G | GS | MP | PER | TS% | 3PAr | FTr | ORB% | DRB% | TRB% | AST% | STL% | BLK% | TOV% |
|--------------|------|------|--------|------|-----|------|-----|------|------|------|------|------|------|------|
| 21679 | 29.0 | 0.0 | 342.0 | 15.9 | 0.6 | 0.0 | 0.5 | 10.3 | 16.6 | 13.4 | 5.2 | 2.0 | 3.5 | 15.6 |
| 21680 | 52.0 | 5.0 | 713.0 | 13.4 | 0.5 | 0.0 | 0.6 | 10.6 | 21.2 | 15.7 | 8.3 | 1.3 | 3.1 | 13.1 |
| 21681 | 64.0 | 64.0 | 2307.0 | 13.0 | 0.5 | 0.3 | 0.2 | 1.4 | 10.3 | 5.8 | 14.6 | 0.9 | 0.4 | 12.1 |
| 21682 | 3.0 | 0.0 | 9.0 | 15.3 | 0.6 | 0.5 | 0.0 | 0.0 | 12.1 | 6.2 | 19.4 | 0.0 | 0.0 | 0.0 |
| 21683 | 45.0 | 0.0 | 388.0 | 11.1 | 0.6 | 0.0 | 0.2 | 8.7 | 26.7 | 17.7 | 3.4 | 0.7 | 4.6 | 20.6 |

```
# Create an empty list that will eventually hold all created feature columns.
feature_columns = []
```

```
resolution_in_Zs = 0.3 # 3/10 of a standard deviation.
```

```
# Create a bucket feature column for offensive win shares.
```

```
OWS_as_a_numeric_column = tf.feature_column.numeric_column("OWS")
```

```
OWS_boundaries = list(np.arange(int(min(train_df_norm['OWS'])),
                                int(max(train_df_norm['OWS'])),
                                resolution_in_Zs))
```

```
OWS = tf.feature_column.bucketized_column(OWS_as_a_numeric_column, OWS_boundaries)
```

```
#Create a bucket feature column for defensive win shares.
```

```
DWS_as_a_numeric_column = tf.feature_column.numeric_column("DWS")
```

```
DWS_boundaries = list(np.arange(int(min(train_df_norm['DWS'])),
                                int(max(train_df_norm['DWS'])),
                                resolution_in_Zs))
```

```
DWS = tf.feature_column.bucketized_column(DWS_as_a_numeric_column,
                                           DWS_boundaries)
```

```
# Create a feature cross of OWS and DWS.
```

```
OWS_x_DWS = tf.feature_column.crossed_column([OWS,DWS],hash_bucket_size=30)
```

```
crossed_feature_1 = tf.feature_column.indicator_column(OWS_x_DWS)
```

```
feature_columns.append(crossed_feature_1)
```

```
# Create a bucket feature column for offensive rebounds.
```

```
ORB_as_a_numeric_column = tf.feature_column.numeric_column("ORB")
```

```
ORB_boundaries = list(np.arange(int(min(train_df_norm['ORB'])),
                                int(max(train_df_norm['ORB'])),
                                resolution_in_Zs))

ORB = tf.feature_column.bucketized_column(ORB_as_a_numeric_column, ORB_boundaries)

#Create a bucket feature column for defensive rebounds.

DRB_as_a_numeric_column = tf.feature_column.numeric_column("DRB")

DRB_boundaries = list(np.arange(int(min(train_df_norm['DRB'])),
                                int(max(train_df_norm['DRB'])),
                                resolution_in_Zs))

DRB = tf.feature_column.bucketized_column(DRB_as_a_numeric_column,
                                          DRB_boundaries)

# Create a feature cross of ORB and DRB.
ORB_x_DRB = tf.feature_column.crossed_column([ORB,DRB],hash_bucket_size=30)
crossed_feature_2 = tf.feature_column.indicator_column(ORB_x_DRB)
feature_columns.append(crossed_feature_2)

# Represent Turnovers as a floating-point value.
Turnovers = tf.feature_column.numeric_column("TOV")
feature_columns.append(Turnovers)

# Represent Steals as a floating-point value.
Steals = tf.feature_column.numeric_column("STL")
feature_columns.append(Steals)

# Represent population as a floating-point value.
Three_Point_Attempts = tf.feature_column.numeric_column("3PA")
feature_columns.append(Three_Point_Attempts)

# Represent population as a floating-point value.
Minutes_Played = tf.feature_column.numeric_column("MP")
feature_columns.append(Minutes_Played)

# Represent population as a floating-point value.
Two_Point_Attempts = tf.feature_column.numeric_column("2PA")
feature_columns.append(Two_Point_Attempts)

# Represent Field Goal Attempts as a floating-point value.
Field_Goal_Attempts = tf.feature_column.numeric_column("FGA")
feature_columns.append(Field_Goal_Attempts)

my_feature_layer = tf.keras.layers.DenseFeatures(feature_columns)
```

```
def plot_the_loss_curve(epochs, mse):
    """Plot a curve of loss vs. epoch."""

    plt.figure()
    plt.xlabel("Epoch")
    plt.ylabel("Mean Squared Error")

    plt.plot(epochs, mse, label="Loss")
    plt.legend()
    plt.ylim([mse.min()*0.95, mse.max() * 1.03])
    plt.show()

print("Defined the plot_the_loss_curve function.")
```

↳ Defined the plot_the_loss_curve function.

```
def train_model(model, dataset, epochs, label_name,
                batch_size=None):

    # Split the dataset into features and label.
    features = {name:np.array(value) for name, value in dataset.items()}
    label = np.array(features.pop(label_name))
    history = model.fit(x=features, y=label, batch_size=batch_size,
                        epochs=epochs, shuffle=True)

    #store epochs separately from the rest of history
    epochs = history.epoch

    # To track the progression of training, gather a snapshot
    # of the model's mean squared error at each epoch.
    hist = pd.DataFrame(history.history)
    mse = hist["mean_squared_error"]

    return epochs, mse

def create_model(my_learning_rate, my_feature_layer):
    #Create and compile a simple linear model
    model = tf.keras.models.Sequential()

    model.add(my_feature_layer)

    # Define the first hidden layer with 32 nodes.
    model.add(tf.keras.layers.Dense(units=32,
                                     activation='relu',
                                     name='Hidden1'))

    # Define the second hidden layer with 20 nodes.
    model.add(tf.keras.layers.Dense(units=20,
                                     activation='relu',
                                     name='Hidden2'))
```

```
# Define the third hidden layer with 10 nodes.
model.add(tf.keras.layers.Dense(units=10,
                                activation='relu',
                                name='Hidden3'))

# Define the output layer.
model.add(tf.keras.layers.Dense(units=1,
                                name='Output'))

model.compile(optimizer=tf.keras.optimizers.Adam(lr=my_learning_rate),
              loss="mean_squared_error",
              metrics=[tf.keras.metrics.MeanSquaredError()])

return model

# The following variables are the hyperparameters.
learning_rate = 0.01
epochs = 40
batch_size = 1000

# Specify the label
label_name = "PTS"

# Establish the model's topography.
my_model = create_model(learning_rate, my_feature_layer)

# Train the model on the normalized training set. We're passing the entire
# normalized training set, but the model will only use the features
# defined by the feature_layer.
epochs, mse = train_model(my_model, train_df_norm, epochs,
                          label_name, batch_size)
plot_the_loss_curve(epochs, mse)


# After building a model against the training set, test that model
# against the test set.
test_features = {name:np.array(value) for name, value in test_df_norm.items()}
test_label = np.array(test_features.pop(label_name)) # isolate the label
print("\n Evaluate the new model against the test set:")
my_model.evaluate(x = test_features, y = test_label, batch_size=batch_size)
```



```
Epoch 1/40
2/2 [=====] - 0s 4ms/step - loss: 0.9237 - mean_squared_error:
Epoch 2/40
2/2 [=====] - 0s 4ms/step - loss: 0.5695 - mean_squared_error:
Epoch 3/40
2/2 [=====] - 0s 3ms/step - loss: 0.2429 - mean_squared_error:
Epoch 4/40
2/2 [=====] - 0s 4ms/step - loss: 0.1132 - mean_squared_error:
Epoch 5/40
2/2 [=====] - 0s 3ms/step - loss: 0.1652 - mean_squared_error:
Epoch 6/40
2/2 [=====] - 0s 3ms/step - loss: 0.1463 - mean_squared_error:
Epoch 7/40
2/2 [=====] - 0s 2ms/step - loss: 0.0737 - mean_squared_error:
Epoch 8/40
2/2 [=====] - 0s 2ms/step - loss: 0.0526 - mean_squared_error:
Epoch 9/40
2/2 [=====] - 0s 3ms/step - loss: 0.0636 - mean_squared_error:
Epoch 10/40
2/2 [=====] - 0s 3ms/step - loss: 0.0664 - mean_squared_error:
Epoch 11/40
2/2 [=====] - 0s 2ms/step - loss: 0.0487 - mean_squared_error:
Epoch 12/40
2/2 [=====] - 0s 3ms/step - loss: 0.0325 - mean_squared_error:
Epoch 13/40
2/2 [=====] - 0s 3ms/step - loss: 0.0357 - mean_squared_error:
Epoch 14/40
2/2 [=====] - 0s 3ms/step - loss: 0.0413 - mean_squared_error:
Epoch 15/40
2/2 [=====] - 0s 3ms/step - loss: 0.0306 - mean_squared_error:
Epoch 16/40
2/2 [=====] - 0s 3ms/step - loss: 0.0243 - mean_squared_error:
Epoch 17/40
2/2 [=====] - 0s 4ms/step - loss: 0.0277 - mean_squared_error:
Epoch 18/40
2/2 [=====] - 0s 4ms/step - loss: 0.0284 - mean_squared_error:
Epoch 19/40
2/2 [=====] - 0s 4ms/step - loss: 0.0233 - mean_squared_error:
Epoch 20/40
2/2 [=====] - 0s 4ms/step - loss: 0.0196 - mean_squared_error:
Epoch 21/40
2/2 [=====] - 0s 3ms/step - loss: 0.0211 - mean_squared_error:
Epoch 22/40
2/2 [=====] - 0s 3ms/step - loss: 0.0213 - mean_squared_error:
Epoch 23/40
2/2 [=====] - 0s 3ms/step - loss: 0.0185 - mean_squared_error:
Epoch 24/40
2/2 [=====] - 0s 4ms/step - loss: 0.0176 - mean_squared_error:
Epoch 25/40
2/2 [=====] - 0s 3ms/step - loss: 0.0183 - mean_squared_error:
Epoch 26/40
2/2 [=====] - 0s 3ms/step - loss: 0.0179 - mean_squared_error:
Epoch 27/40
2/2 [=====] - 0s 3ms/step - loss: 0.0163 - mean_squared_error:
Epoch 28/40
2/2 [=====] - 0s 3ms/step - loss: 0.0159 - mean_squared_error:
Epoch 29/40
```



```
2/2 [=====] - 0s 3ms/step - loss: 0.0163 - mean_squared_error:
Epoch 30/40
2/2 [=====] - 0s 4ms/step - loss: 0.0153 - mean_squared_error:
Epoch 31/40
2/2 [=====] - 0s 4ms/step - loss: 0.0148 - mean_squared_error:
Epoch 32/40
2/2 [=====] - 0s 4ms/step - loss: 0.0153 - mean_squared_error:
Epoch 33/40
2/2 [=====] - 0s 4ms/step - loss: 0.0147 - mean_squared_error:
Epoch 34/40
2/2 [=====] - 0s 4ms/step - loss: 0.0142 - mean_squared_error:
Epoch 35/40
2/2 [=====] - 0s 4ms/step - loss: 0.0141 - mean_squared_error:
Epoch 36/40
2/2 [=====] - 0s 3ms/step - loss: 0.0138 - mean_squared_error:
Epoch 37/40
2/2 [=====] - 0s 4ms/step - loss: 0.0136 - mean_squared_error:
Epoch 38/40
2/2 [=====] - 0s 3ms/step - loss: 0.0135 - mean_squared_error:
Epoch 39/40
2/2 [=====] - 0s 3ms/step - loss: 0.0131 - mean_squared_error:
Epoch 40/40
2/2 [=====] - 0s 4ms/step - loss: 0.0129 - mean_squared_error:
```



Evaluate the new model against the test set:

```
1/1 [=====] - 0s 1ms/step - loss: 0.0261 - mean_squared_error:
[0.02606414072215557, 0.02606414072215557]
```